

The Fortress Language Specification

May 28, 2012

Contents

1	Introduction	1
2	Overview of Fortress	2
3	Getting Started	3
4	Programs	4
5	Types	5
5.1	Kinds of Types	5
5.2	Relationships between Types	6
5.3	Trait Types	7
5.3.1	Object Trait Types	7
5.4	Object Expression Types	7
5.5	Tuple Types	7
5.6	Arrow Types	8
5.7	Function Types	9
5.8	Special Types	9
5.9	Types in the Fortress Standard Libraries	9
5.10	Intersection and Union Types	10
5.11	Type Aliases	11
6	Lexical Structure	12
6.1	Reserved Words	12
7	Names and Declarations	13
8	Variables	14
9	Functions	15
10	Traits	16
11	Objects	17
12	Expressions	18
12.1	Literals	18
12.2	Identifier References	18
12.3	Dotted Field Accesses	18
12.4	Function Calls	18
12.5	Operator Applications	18
12.6	Tuple Expressions	18

12.7	Aggregate Expressions	18
12.8	Function Expressions	18
12.9	Blocks	18
12.10	Do and Do-Also Expressions	19
12.11	Label and Exit	20
12.12	While Loops	20
12.13	Generators	22
12.14	For Expressions	22
12.15	Ranges	22
12.16	Reduction Expressions	22
12.17	Comprehensions	22
12.18	If Expressions	22
12.19	Case Expressions	23
12.20	Typecase Expressions	23
12.21	Atomic Expressions	23
12.22	Throw Expressions	23
12.23	Try Expressions	23
12.24	Type Ascription	23
12.25	Asif Expressions	23
13	Exceptions	24
14	Operators	25
14.1	Operator Names	25
14.2	Operator Names	25
14.3	Operator Precedence and Associativity	26
14.4	Operator Fixity	28
14.5	Chained and Multifix Operators	28
14.6	Enclosing Operators	29
14.7	Conditional Operators	30
14.8	Big Operators	30
14.9	Juxtaposition	31
14.10	Operator Declarations	33
14.10.1	Multifix Operator Declarations	33
14.10.2	Infix Operator Declarations	34
14.10.3	Prefix Operator Declarations	34
14.10.4	Postfix Operator Declarations	34
14.10.5	Nofix Operator Declarations	34
14.10.6	Bracketing Operator Declarations	35
14.11	Overview of Operators in the Fortress Standard Libraries	35
14.11.1	Prefix Operators	35
14.11.2	Postfix Operators	35
14.11.3	Enclosing Operators	35
14.11.4	Exponentiation	35
14.11.5	Superscript Operators	36
14.11.6	Subscript Operators	36
14.11.7	Multiplication, Division, Modulo, and Remainder Operators	36
14.11.8	Addition and Subtraction Operators	37
14.11.9	Intersection, Union, and Set Difference Operators	37
14.11.10	Minimum and Maximum Operators	37
14.11.11	GCD, LCM, and CHOOSE Operators	37
14.11.12	Comparisons Operators	37

14.11.13	Logical Operators	38
14.11.14	Conditional Operators	38
15	Overloading and Dispatch	39
15.1	Principles of Overloading	39
15.2	Visibility to Named Procedure Calls	40
15.3	Applicability to Named Procedure Calls	40
15.4	Applicability to Dotted Method Calls	41
15.5	Applicability for Procedures with Varargs and Keyword Parameters	41
15.6	Overloading Resolution	41
16	Coercion	42
17	Type Inference	43
18	Components and APIs	44
A	Simplified Grammar for Application Programmers and Library Writers	45
A.1	Components and API	45
A.2	Top-level Declarations	46
A.3	Trait Declaration	46
A.4	Object declaration	46
A.5	Variable Declaration	47
A.6	Function Declaration	47
A.7	Headers	48
A.8	Parameters	49
A.9	Method Declaration	49
A.10	Method Parameters	49
A.11	Field Declarations	50
A.12	Abstract Filed Declaration	50
A.13	Expressions	50
A.14	Expressions Enclosed by Keywords	51
A.15	Local Declarations	53
A.16	Literals	53
A.17	Types	53
A.18	Symbols and Operators	54
A.19	Identifiers	54

Chapter 1

Introduction

Chapter 2

Overview of Fortress

Chapter 3

Getting Started

Chapter 4

Programs

Chapter 5

Types

This chapter should be revised according to Victor’s tuple story (his email titled “Tuple story” on 12/07/07).

No dimensions and units.

The Fortress type system supports a mixture of nominal and structural types. Specifically, Fortress provides nominal trait types and some constructs that combine types structurally. There are also a few special types. Most trait types are defined by trait declarations in a program (often in libraries), but a few are built-in.

Victor: This doesn’t cover the special types, dimension types, or object expression types. But I think it’s okay for this intro, as long as we don’t assert that these are all the types.

Every expression has a *static type*, and every value has an *ilk* (also known as a *dynamic or runtime type*). Fortress programs are checked before they are executed to ensure that if an expression e evaluates to a value v , the runtime type of v is a subtype of the static type of e . Sometimes we abuse terminology by saying that an expression has the runtime type of the value it evaluates to (see Section ?? for a discussion about evaluation of expressions).

Some types may be parameterized by *static parameters*; we call these types *generic types*. See Chapter ?? for a discussion of static parameters. A static parameter may be instantiated with a type or a value depending on whether it is a type parameter.

Victor: I don’t know if the follow really belongs here. It isn’t complete in any case, because of union/intersection types.

Two types are identical if and only if they are the same kind and their names and static arguments (if any) are identical. Types are related by several relationships as described in Section 5.2.

Syntactically, the positions in which a type may legally appear (i.e., *type contexts*) is determined by the nonterminal *Type* in the Fortress grammar, defined in Appendix ??.

5.1 Kinds of Types

Fortress supports the following kinds of types:

- trait types,
- object expression types,
- tuple types,
- arrow types,
- function types, and
- three special types: `Any`, `BottomType` and `()`.

Object expression types, function types and `BottomType` are not first-class types: they cannot be written in a program. However, some values have object expression types and function types and `throw` and `exit` expressions have the type `BottomType`.

Collectively, trait types and object expression types are *defined* types; the trait and object declarations and object expressions are the types' *definitions*.

Victor:

I don't know if you like this terminology for trait types and object expression types. Putting this all together allowed some repetition to be eliminated, and emphasizes the similarity, which I think you wanted to retain. I still want to separate object expression types from object trait types. I think the difference here is not as evident as it might be if we handled generics properly: it is the type environment of object expression types that make them a pain, and I'm not dealing with that for the most part.

5.2 Relationships between Types

We define two relations on types: *subtype* and *exclusion*.

add coercion in full language

The subtype relation is a partial order on types that defines the type hierarchy; that is, it is reflexive, transitive and antisymmetric. `Any` is the top of the type hierarchy: all types are subtypes of `Any`. `BottomType` is the bottom of the type hierarchy: all types are supertypes of `BottomType`. For types T and U , we write $T \preceq U$ when T is a subtype of U , and $T \prec U$ when $T \preceq U$ and $T \neq U$; in the latter case, we say T is a *strict* subtype of U . We also say that T is a *supertype* of U if U is a subtype of T . Thus, in the execution of a valid Fortress program, an expression's runtime type is always a subtype of its static type. We say that a value is *an instance of* its runtime type and of every supertype of its runtime type; immediate subtypes of `Any` comprises of tuple types, arrow types, `()`, and `Object`.

The exclusion relation is a symmetric relation between two types whose intersection is `BottomType`. Because `BottomType` is uninhabited (i.e., no value has type `BottomType`), types that exclude each other are disjoint: no value can have a type that is a subtype of two types that exclude each other.

The converse is technically not true: we can define two traits with method declarations such that no trait or object can extend them both, so no value can have a type that is a subtype of both trait types, but the trait types do not exclude each other. For example: trait `A f(self, y: Object) = 1; end` trait `B f(x: Object, self) = 2; end` trait `C f() = 3; end`

No pair of the three trait types defined above can be extended but none of them exclude any of the others.

Note that `BottomType` excludes every type including itself (because the intersection of `BottomType` and any type is `BottomType`), and also that no type other than `BottomType` excludes `Any` (because the intersection of any type T and `Any` is T). `BottomType` is the only type that excludes itself. Also note that if two types exclude each other then any subtypes of these types also exclude each other.

Fortress also allows *coercion* between types (see Chapter ??). A coercion from T to U is defined in the declaration of U . We write $T \rightarrow U$ if U defines a coercion from T . We say that T *can be coerced to* U , and write $T \rightsquigarrow U$, if U defines a coercion from T or any supertype of T : $T \rightsquigarrow U \iff \exists T' : T \preceq T' \wedge T' \rightarrow U$.

The Fortress type hierarchy is acyclic with respect to both subtyping and coercion relations except for the following:

- The trait `Any` is a single root of the type hierarchy and it forms a cycle as described in Chapter ??.
- There exists a bidirectional coercion between two tuple types if and only if they have the same sorted form.

These relations are defined more precisely in the following sections describing each kind of type in more detail. Specifically, the relations are the smallest ones that satisfy all the properties given in those sections (and this one).

5.3 Trait Types

Syntax:

$$Type ::= TraitType$$

A *trait type* is a named type defined by a trait or object declaration. It is an *object trait type* if it is defined by an object declaration.

Victor: The following text really belongs to traits, not to types. Traits are declared by trait and object declarations (described in Chapter 10 and Chapter 11). A trait has a *trait type* of the same name.

A trait type is a subtype of every type that appears in the `extends` clause of its definition. In addition, every trait type is a subtype of `Any`.

A trait declaration may also include an `excludes` clause, in which case the defined trait type excludes every type that appears in that clause. Every trait type also excludes every arrow type and every tuple type, and the special types `()` and `BottomType`.

5.3.1 Object Trait Types

An object trait type is a trait type defined by an object declaration. Object declarations do not include `excludes` clauses, but an object trait type cannot be extended. Thus, an object trait type excludes any type that is not its supertype.

5.4 Object Expression Types

An object expression type is defined by an object expression (described in Section ??) and the program location of the object expression. Every evaluation of a given object expression has the same object expression type. Two object expressions at different program locations have different object expression types.

Like trait types, an object expression type is a subtype of all trait types that appear in the `extends` clause of its definition, as well as the type `Any`.

Like object trait types, an object expression type excludes any type that is not its supertype.

5.5 Tuple Types

Syntax:

$$\begin{aligned} TupleType &::= (Type, TypeList) \\ TypeList &::= Type(, Type)^* \end{aligned}$$

Victor: The following two sentences don't belong in this section: A tuple expression is an ordered sequence of expressions separated by commas and enclosed in parentheses. See Section ?? for a discussion of tuple expressions.

A tuple type consists of a parenthesized, comma-separated list of two or more types.

Every tuple type is a subtype of `Any`. No other type encompasses all tuple types. Tuple types cannot be extended by trait types. Tuple types are covariant: a tuple type X is a subtype of tuple type Y if and only if they have the same number of element types and each element type of X is a subtype of the corresponding element type of Y .

A tuple type excludes any non-tuple type other than `Any`. A tuple type excludes every tuple type that does not have the same number of element types. Also, tuple types that have the same number of element types exclude each other if any pair of corresponding element types exclude each other.

Intersection of nonexclusive tuple types are defined elementwise; the intersection of nonexclusive tuple type X and Y is a tuple type with exactly corresponding elements, where the type in each element type is the intersection of the types in the corresponding element types of X and Y . Note that intersection of any exclusive types is `BottomType` as described in Section 5.8.

5.6 Arrow Types

Arrow types should include `io`-ness.

Arrow types with contracts

Function contracts consist of three parts: a `requires` part, an `ensures` part, and an `invariant` part. All three parts are evaluated in the scope of the function body extended with a special variable `arg`, bound to an immutable array of all function arguments. (This array is useful for describing contracts in the presence of higher-order functions; see Section 5.6).

Is the special variable `arg` reserved? What if one of the function’s parameters is named `arg`?

At runtime, when a function is bound to a variable or parameter `f` whose type includes a contract, the contract is evaluated when the function is called through a reference to `f`. (Note that this contract is distinct from the contract attached to the function bound to `f`, which is also evaluated upon a call to `f`). This contract is evaluated in the enclosing scope of the arrow type, extended with any keyword argument names provided in the arrow type, and the special variable `arg` bound to an immutable array containing the arguments provided at the call site (in the order provided at the call site). If the `requires` clauses stipulated in the type of `f` are not satisfied, a `CallerViolation` exception is thrown. Otherwise, if the `requires` clauses attached to the function bound to `f` is not satisfied, a `ContractBinding` exception is thrown. If any other part of the contract of the function bound to `f` is not satisfied, a `CalleeViolation` exception is thrown. Otherwise, if any part of the contract stipulated in the type of `f` is not satisfied, a `ContractBinding` exception is thrown.

Syntax:

```

Type      ::= ArgType → Type Throws?
ArgType   ::= ( (Type,)* Type... )
           | TupleType

```

The static type of an expression that evaluates to a function value is an *arrow type* (or possibly an intersection of arrow types). An arrow type has three constituent parts: a *parameter type*, a *return type*, and a set of *exception types*. (Multiple parameters or return values are handled by having tuple types for the parameter type or return type respectively.)

Victor: I’d rather say “an exception type”, where that type can be a union type, but I’m avoiding mentioning union types for now.

Syntactically, an arrow type consists of the parameter type followed by the token `→`, followed by the return type, and optionally a `throws` clause that specifies the set of exception types.

Victor: What about `io`? Ignoring that for now.

If there is no `throws` clause, the set of exception types is empty.

The parameter type of an arrow type may end with a “varargs entry”, `T...`. The resulting parameter type is not a first-class type; rather, it is the union of all the types that result from replacing the varargs entry with zero or more entries of the type `T`. For example, `(T...)` is the union of `()`, `T`, `(T, T)`, `(T, T, T)`, and so on.

We say that an arrow type is *applicable* to a type `A` if `A` is a subtype of the parameter type of the arrow type. (If the parameter type has a varargs entry, then `A` must be a subtype of one of the types in the union defined by the parameter type.)

Every arrow type is a subtype of `Any`. No other type encompasses all arrow types. A type parameter can be instantiated with an arrow type. Arrow types cannot be extended by trait types. Arrow types are covariant in their return type and exception types and contravariant in their parameter type; that is, arrow type “`A → B throws C`” is a subtype of arrow type “`D → E throws F`” if and only if:

- `D` is a subtype of `A`,
- `B` is a subtype of `E`, and
- for all `X` in `C`, there exists `Y` in `F` such that `X` is a subtype of `Y`.

For arrow types S and T , we say that S is *more specific than* T if the parameter type of S is a subtype of the parameter type of T . We also say that S is *more restricted than* T if the return type of S is a subtype of the return type of T and for every X in the set of exception types of S , there exists Y in the set of exception types of T such that X is a subtype of Y . Thus, S is a subtype of T if and only if T is more specific than S and S is more restricted than T .

An arrow type excludes any non-arrow type other than `Any` and the function types that are its subtypes (see Section 5.7). However, arrow types do not exclude other arrow types because of overloading as described in Chapter ??.

Here are some examples:

5.7 Function Types

Function types are the runtime types of function values. We distinguish them from arrow types to handle overloaded functions. A function type consists of a finite set of arrow types. However, not every set of arrow types is a well-formed function type. Rather, a function type F is *well-formed* if, for every pair (S_1, S_2) of distinct arrow types in F , the following properties hold:

- the parameter types of S_1 and S_2 are not the same,
- if S_1 is more specific than S_2 then S_1 is more restricted than S_2 , and
- if the intersection of the parameter types of S_1 and S_2 is not `BottomType` (i.e., the parameter types of S_1 and S_2 do not exclude each other) then F has some constituent arrow type that is more specific than both S_1 and S_2 (recall that the more specific relation is reflexive, so the required constituent type may be S_1 or S_2).

Henceforth, we consider only well-formed function types. The overloading rules ensure that all function values have well-formed function types.

We say that a function type F is *applicable* to a type A if any of its constituent arrow types is applicable to A . We extend this terminology to values of the corresponding types. That is, for example, we say that a function of type F is applicable to a value of type A if F is applicable to A . Note that if a well-formed function type F is applicable to a type A , then among all constituent arrow types of F that are applicable to A (and there must be at least one), one is more specific than all the others. We say that this constituent type is the *most specific* type of F applicable to A .

A function type is a subtype of each of its constituent arrow types, and also of `Any`. Like object trait types and object expression types, a function type excludes every type that is not its supertype.

5.8 Special Types

Fortress provides three special types: `Any`, `BottomType` and `()`. The type `Any` is the top of the type hierarchy: it is a supertype of every type. The only type it excludes is `BottomType`.

The type `()` is the type of the value `()`. Its only supertype (other than itself) is `Any`, and it excludes every other type.

Fortress provides a special *bottom type*, `BottomType`, which is an uninhabited type. No value in Fortress has `BottomType`; `throw` and `exit` expressions have `BottomType`. `BottomType` is a subtype of every type and it excludes every type (including itself). As mentioned above, `BottomType` is not a first-class type: programmers must not write `BottomType`.

Victor's comments The complex numbers have precision like the integers, but are they intended to be only integral complex numbers? That seems odd to me. Also, why say “imaginary and complex”, since we don't provide a separate type for just the imaginary numbers? Guy's going to provide APIs for complex numbers.

5.9 Types in the Fortress Standard Libraries

The Fortress standard libraries define simple standard types for literals such as `BooleanLiteral[b]`, `()` (pronounced “void”), `Character`, `String`, and `Numeral[n, m, r, v]` for appropriate values of b , n , m , r , and v (See Section 12.1

for a discussion of Fortress literals). Moreover, there are several simple standard numeric types. These types are mutually exclusive; no value has more than one of them. Values of these types are immutable.

The numeric types share the common supertype `Number`. Fortress includes types for arbitrary-precision integers (of type \mathbb{Z}), their unsigned equivalents (of type \mathbb{N}), rational numbers (of type \mathbb{Q}), real numbers (of type \mathbb{R}), complex numbers (of type \mathbb{C}), fixed-size representations for integers including the types $\mathbb{Z}8$, $\mathbb{Z}16$, $\mathbb{Z}32$, $\mathbb{Z}64$, $\mathbb{Z}128$, their unsigned equivalents $\mathbb{N}8$, $\mathbb{N}16$, $\mathbb{N}32$, $\mathbb{N}64$, $\mathbb{N}128$, floating-point numbers (described below), intervals (of type `Interval[X]`, abbreviated as $\langle X \rangle$, where X can be instantiated with any number type), and imaginary and complex numbers of fixed size (in rectangular form with types $\mathbb{C}n$ for $n = 16, 32, 64, 128, 256$ and polar form with type `Polar[X]` where X can be instantiated with any real number type).

The Fortress standard libraries also define other simple standard types such as `Any`, `Object`, `Exception`, `Boolean`, and `BooleanInterval` as well as low-level binary data types such as `LinearSequence`, `HeapSequence`, and `BinaryWord`. See Parts ?? and ?? for discussions of the Fortress standard libraries.

5.10 Intersection and Union Types

For every finite set of types, there is a type denoting a unique *intersection* of those types. The intersection of a set of types S is a subtype of every type $T \in S$ and of the intersection of every subset of S . There is also a type denoting a unique *union* of those types. The union of a set of types S is a supertype of every type $T \in S$ and of the union of every subset of S . Neither intersection types nor union types are first-class types; they are used solely for type inference (as described in Chapter 17) and they cannot be expressed directly in programs.

The intersection of a set of types S is equal to a named type U when any subtype of every type $T \in S$ and of the intersection of every subset of S is a subtype of U . Similarly, the union of a set of types S is equal to a named type U when any supertype of every type $T \in S$ and of the union of every subset of S is a supertype of U . For example:

```
trait S comprises {U, V} end
trait T comprises {V, W} end
trait U extends S excludes W end
trait V extends {S, T} end
trait W extends T end
```

because of the `comprises` clauses of S and T and the `excludes` clause of U , any subtype of both S and T must be a subtype of V . Thus, $V = S \cap T$.

Intersection types (denoted by \cap) possess the following properties:

- Commutativity: $T \cap U = U \cap T$.
- Associativity: $S \cap (T \cap U) = (S \cap T) \cap U$.
- Subsumption: If $S \preceq T$ then $S \cap T = S$.
- Preservation of shared subtypes: If $T \preceq S$ and $T \preceq U$ then $T \preceq S \cap U$.
- Preservation of supertype: If $S \preceq T$ then $\forall U. S \cap U \preceq T$.
- Distribution over union types: $S \cap (T \cup U) = (S \cap T) \cup (S \cap U)$.

Union types (denoted by \cup) possess the following properties:

- Commutativity: $T \cup U = U \cup T$.
- Associativity: $S \cup (T \cup U) = (S \cup T) \cup U$.
- Subsumption: If $S \preceq T$ then $S \cup T = T$.
- Preservation of shared supertypes: If $S \preceq T$ and $U \preceq T$ then $S \cup U \preceq T$.
- Preservation of subtype: If $T \preceq S$ then $\forall U. T \preceq S \cup U$.
- Distribution over intersection types: $S \cup (T \cap U) = (S \cup T) \cap (S \cup U)$.

5.11 Type Aliases

Syntax:

$$TypeAlias ::= \text{type } Id \text{ StaticParams? } = Type$$

Fortress allows names to serve as aliases for more complex type instantiations. A *type alias* begins with `type` followed by the name of the alias type, followed by optional static parameters, followed by `=`, followed by the type it stands for. Parameterized type aliases are allowed but recursively defined type aliases are not. Here are some examples:

```
type IntList = List[Z64]
type BinOp = Float × Float → Float
type SimpleFloat[nat e, nat s] = DetailedFloat[Unity, e, s, false, false, false, false, true]
```

All uses of type aliases are expanded before type checking. Type aliases do not define new types nor nominal equivalence relations among types.

Chapter 6

Lexical Structure

6.1 Reserved Words

The following tokens are *reserved words*:

BIG	FORALL	SI_unit	absorbs	abstract	also	api
asif	at	atomic	bool	case	catch	coerce
coerces	component	comprises	default	dim	do	elif
else	end	ensures	except	excludes	exit	export
extends	finally	fn	for	forbid	from	getter
hidden	if	import	int	invariant	io	juxtaposition
label	most	nat	native	object	of	opr
or	outcome	override	private	property	provided	requires
self	settable	setter	spawn	syntax	test	then
throw	throws	trait	try	tryatomic	type	typecase
typed	unit	value	var	where	while	widens
with	wrapped					

Victor: I don't think 'or' should be reserved. It is only so for its occurrence in 'widens or coerces' but we can recognize it specially in that context, which is never ambiguous because 'widens' and 'coerces' are reserved.

The following operators on units are also reserved words:

cubed cubic in inverse per square squared

To avoid confusion, Fortress reserves the following tokens:

goto idiom public pure reciprocal static

They do not have any special meanings but they cannot be used as identifiers.

Victor: Some other words we might want to reserve: subtype, subtypes, is, coercion, function, exception, match

Chapter 7

Names and Declarations

Chapter 8

Variables

Chapter 9

Functions

Chapter 10

Traits

```
export Executable  
run() = println("Hello, world!")
```

File	::=	CompilationUnit
		Imports [?] Exports Decls [?]
		Imports [?] AbsDecls
		Imports AbsDecls [?]
CompilationUnit	::=	Component
		Api

Chapter 11

Objects

Chapter 12

Expressions

12.1 Literals

A literal (Section ??) denotes a fixed, unchanging value.

Literal	::=	()
		BooleanLiteral
		CharacterLiteral
		StringLiteral
		NumericLiteral

The type of the literal `()` is `()`. The type of a Boolean literal is `Boolean`. The type of a character literal is `Character`. The type of a string literal is `String`. The type of a numeric literal is `ℚ` if it contains a `'.'` character, and otherwise is `ℕ`.

Evaluation of a literal always completes normally and produces the value represented by the literal.

All literals represent value objects; therefore their values do not have object identity.

12.2 Identifier References

12.3 Dotted Field Accesses

12.4 Function Calls

12.5 Operator Applications

12.6 Tuple Expressions

12.7 Aggregate Expressions

12.8 Function Expressions

12.9 Blocks

A block is a series of one or more declarations and expressions that are to be evaluated sequentially. A block is not in itself syntactically an expression, but appears as part of a `do` expression Section 12.10, `label` expression Sec-

tion 12.11, while expression Section 12.12, for expression Section 12.14, if expression Section 12.18, case expression Section 12.19, typecase expression Section 12.20, or try expression Section 12.23. For purposes of type-checking and evaluation, it may be regarded as an expression.

```
Block          ::= BlockElement+
BlockElement   ::= LocalVarFnDecl
                | Expr ( , GeneratorList )?
```

It is a static error if the last BlockElement in a Block is an expression of the form “ $a = b$ ” that cannot be regarded as a local declaration. It is a static error if any block element other than the last one is an expression whose type is not $()$. If the last block element of a block is a declaration, then the type of the block is $()$; if the last block element of a block is an expression, then the type of the block is the type of that expression.

Every block introduces a new scope. For the scoping behavior of local declarations within a block, see Section ?? and Section ??.

Not shown in the grammar is the fact that adjacent block elements may be separated by a semicolon ‘;’ if desired. This is necessary if the block elements appear in the same line of source code. If a newline occurs between two adjacent block elements, then no semicolon is necessary.

A block is evaluated by evaluating its block elements sequentially, in order from left to right; evaluation of each block element must complete before evaluation of the next can begin. If evaluation of any block element completes abruptly for some reason, then evaluation of the block completes abruptly for the same reason, and no further block elements are evaluated. If the evaluations of all block elements complete normally, then the value of the block is the value of the last block element.

If it is desired to have, as a non-final block element, an expression e whose type is not $()$, one may instead write the local variable declaration $_ = e$, thus signifying explicitly that the value of e is to be discarded (by binding an anonymous variable to the value).

If it is desired to have, as a final block element, an equality test expression of the form $a = b$, simply enclose the expression in parentheses.

Here is an example of a `do` expression that contains a single block having six block elements:

```
do
  f(w: ℤ32) = w + 1    * Local function declaration
  y = x + 1            * Local variable declaration (immutable)
  println(“y is” y)    * Expression (has a useful side effect)
  var z: ℝ64 = 0        * Local variable declaration (mutable)
  z += f(y)            * Compound assignment
  |z|                  * Expression (the value of the block)
end
```

12.10 Do and Do-Also Expressions

In the simplest case, the keywords `do` and `end` surround a single block to be executed. In the more general case, a `do` expression can contain multiple blocks to be executed independently (perhaps but not necessarily, concurrently). Each block may be governed by its own `atomic` modifier.

```
DoExpr          ::= ( DoFront also )* DoFront end
DoFront          ::= ( at Expr )? atomic? do Block?
```

If a `do` expression contains a single Block, then the type of the `do` expression is the type of the block. If a `do` expression contains two or more blocks (separated by `also`), then the type of the `do` expression is $()$, and it is a static error if any of the blocks has a type that is not $()$.

If the `do` keyword preceding any block of a multi-block `do` expression is preceded by an `atomic` modifier, it is as if (a) the following block were enclosed by a new pair of `do` and `end` keywords to form a single-block `do` expression, and (b) the `atomic` modifier appeared instead before that single-block `do` expression, thus forming an atomic expression. Thus, allowing `atomic` to appear as part of a multi-block `do` expression is merely a syntactic convenience that allows programs to be slightly shorter and perhaps read more naturally. For example:

```
atomic do
  x += 1
  y += 2
also atomic do
  b += 1
  y += 3
end
```

is simply an abbreviation of

```
do
  atomic do
    x += 1
    y += 2
  end
also do
  atomic do
    b += 1
    y += 3
  end
end
```

A `do` expression with a single block is evaluated by evaluating the block. If the block completes abruptly for some reason, then the `do` expression completes abruptly for the same reason. Otherwise, the value of the `do` expression is the value of the block.

A `do` expression with multiple blocks is evaluated by evaluating all the blocks independently (perhaps, but not necessarily, concurrently). If more blocks complete abruptly for some reason, then the `do` expression completes abruptly for one of those reasons (evaluation of other blocks may or may not be completed, but evaluation of the `do` expression is not complete until the evaluations of all blocks have been terminated). Otherwise, evaluation of the `do` expression completes only after evaluation of all blocks is complete, and the value of the `do` expression is `()`.

12.11 Label and Exit

12.12 While Loops

The `while` statement evaluates a test and a `do` expression, sequentially and repeatedly, until the test fails. The test may be either a Boolean expression or a generator binding.

```
WhileExpr      ::= while Generator DoExpr
Generator      ::= GeneratorBinding
                | Expr
GeneratorBinding ::= Id ← Expr
                | ( Id , IdList ) ← Expr
```


If the Generator is an expression, it is a static error if the type of the expression does not conform to `Boolean`. If the Generator is a generator binding, it is a static error if the type of the expression in the generator binding does not conform to `Condition[[T]]` for some `T`. The type of a `while` expression is `()`.

A `while` expression with an expression is evaluated by first evaluating `Expr`. If this evaluation completes abruptly for some reason, evaluation of the `while` expression completes abruptly for the same reason. Otherwise, evaluation continues by making a choice based on the resulting value `v`. If `v` is `false`, no further action is taken; evaluation of the `while` expression completes normally with value `()`. But if `v` is `true`, then the `do` expression is evaluated. If this evaluation completes abruptly for some reason, evaluation of the `while` expression completes abruptly for the same reason; otherwise, the entire `while` expression is evaluated again (beginning by re-evaluating `Expr`). For example:

```
do
  * Print the numbers 0 through 9 on separate lines.
  k: ℤ := 0
  while (k < 10) do
    println(k)
    k += 1
  end
end
```

A `while` expression with a generator binding is evaluated by first evaluating the `Expr` in the generator binding. If this evaluation completes abruptly for some reason, evaluation of the `while` expression completes abruptly for the same reason. Otherwise, evaluation continues by making a choice based on the resulting value `v`. If `v` does not contain a value, no further action is taken; evaluation of the `while` expression completes normally with value `()`. If `v` contains a value `w`, then the pattern in the generator binding is matched to that value `w` and then the `do` expression is evaluated, and variables bound by the pattern are visible within the `do` expression. If this evaluation completes abruptly for some reason, evaluation of the `while` expression completes abruptly for the same reason; otherwise, the entire `while` expression is evaluated again (beginning by re-evaluating `Expr`). For example:

```
object Chain(item: ℤ, link: Maybe[[Chain]]) end

printValues(x: Maybe[[Chain]]): () = do
  cursor: Maybe[[Chain]] = x
  while (next ← cursor) do
    println(next.item)
    cursor := next.link
  end
end

run() = printValues(Just Chain(1, Just Chain(2, Just Chain(3, Just Chain(4, None))))))
```

12.13 Generators

12.14 For Expressions

12.15 Ranges

12.16 Reduction Expressions

12.17 Comprehensions

12.18 If Expressions

The `if` statement allows conditional evaluation of a block or conditional evaluation of at most one of a set of blocks. The choice is made by sequentially executing one or more tests and choosing the first statement for which a test succeeds. Each test may be either a Boolean expression or a generator binding.

```
IfExpr      ::=  if Generator then Block ElifClause* ElseClause? end
              |  [(] if Generator then Block ElifClause* ElseClause end? [)]
Generator    ::=  GeneratorBinding
              |  Expr
GeneratorBinding ::=  Id ← Expr
              |  ( Id , IdList ) ← Expr
ElifClause   ::=  elif Expr then Block
ElseClause   ::=  else Block
```

An `if` expression consists of `if` followed by a Generator clause (discussed in Section 12.13), followed by `then`, a Block, a possibly empty sequence of `elif` clauses (each consisting of `elif` followed by a Generator clause, `then`, and a Block), an optional `else` clause (consisting of `else` followed by a Block), and finally `end`. The reserved word `end` may be elided if the `if` expression is immediately enclosed by parentheses; in such a case, the `else` clause is required, not optional.

Each Block is a series of one or more block elements (declarations and expressions). See Section ?? for a description of the various syntactic and semantic properties of blocks.

For each Generator, if it is an expression, it is a static error if the type of the expression does not conform to Boolean; if it is a generator binding, it is a static error if the type of the expression in the generator binding does not conform to $\text{Condition}[T]$ for some T . If the `if` expression has no `else` clause, it is a static error if the type of the block after the first `then`, or in any `elif` clause, is not $()$. The type of an `if` expression is the union of the types of all its blocks.

An `if` expression whose first Generator is an expression is evaluated by first evaluating `Expr`. If this evaluation completes abruptly for some reason, evaluation of the `if` expression completes abruptly for the same reason; otherwise, evaluation continues by making a choice based on the resulting value v . If v is *true*, then the first block is evaluated. If this evaluation completes abruptly for some reason, evaluation of the `if` expression completes abruptly for the same reason; otherwise, the value of the `if` expression is value of this block. If v is *false*, then `elif` clauses are considered (see below).

An `if` expression whose first Generator is a generator binding is evaluated by first evaluating the `Expr` in the generator binding. If this evaluation completes abruptly for some reason, evaluation of the `if` expression completes abruptly for the same reason; otherwise, evaluation continues by making a choice based on the resulting value v . If v contains a value w , then the pattern in the generator binding is matched to that value w and then the first block is evaluated and variables bound by the pattern are visible within the block. If this evaluation completes abruptly for some reason, evaluation of the `if` expression completes abruptly for the same reason; otherwise, the value of the `if` expression is value of this block. If v does not contain a value, then `elif` clauses are considered (see below).

If evaluation of an `if` expression must consider `elif` clauses, they are examined sequentially, working from left to right, treating each `??` and each block in exactly the same manner as the first generator and first block of the `if` expression. As soon as a generator is found that produces the value *true* or a condition value *v* that contains another value *w*, the corresponding block is evaluated, and its value becomes the value of the `if` expression; but if evaluation of any Generator or block completes abruptly for some reason, evaluation of the `if` expression completes abruptly for the same reason. If consideration of `elif` clauses does not result in evaluating an block (possibly because no `elif` clauses are present), then any `else` clause is considered.

If evaluation of an `if` expression must consider any `else` clause, there are two cases. If an `else` clause is present, then its block is evaluated, and its value becomes the value of the `if` expression; but if evaluation of the block completes abruptly for some reason, evaluation of the `if` expression completes abruptly for the same reason. If no `else` clause is present, then evaluation of the `if` expression completes normally with value `()`.

Examples:

```
if x > 0 then println x end
```

```
if x > 0 then
  println(x "is positive")
else
  println(x "is nonpositive")
end
```

```
println(x "is" (if v > 0 then "positive" else "nonpositive"))
```

```
z = if x < 0 then 0
    elif x ∈ {1, 2, 3} then 3
    elif x ∈ {4, 5, 6} then 6
    else 9 end
```

12.19 Case Expressions

12.20 Typecase Expressions

12.21 Atomic Expressions

12.22 Throw Expressions

12.23 Try Expressions

12.24 Type Ascription

12.25 Asif Expressions

Chapter 13

Exceptions

Chapter 14

Operators

Multifix operators and dimensions and units are not yet supported.

Synonym Operators: Victor's email titled "Synonyms [Fwd: Re: What the heck are 0-width spaces for?]" on 09/12/07

Operators are like functions or methods; operator declarations are described in Chapter ?? and operator applications are described in Section ??, Section ??, and Section 12.17. Operators are not values; when an operator is passed as an argument to a function, it should be eta expanded. Just as functions or methods may be overloaded (see Chapter ?? for a discussion of overloading), so operators may have overloaded declarations of the same fixity. Operator declarations with the same operator name but with different fixities are valid declarations because it is always unambiguous which declaration shall be applied to an application of the operator. Calls to overloaded operators are resolved first via the fixity of the operators based on the context of the calls. Then, among the applicable declarations with that fixity, the most specific declaration is chosen.

Most operators can be used as prefix, infix, postfix, or nofix operators as described in Section 14.4 (nofix operators take no arguments); the fixity of an operator is determined syntactically, and the same operator may have declarations for multiple fixities. A simple example is that ‘`-`’ may be either infix or prefix, as is conventional. As another example, ‘`!`’ may be a postfix operator that computes factorial when applied to integers. These operators might not be used as enclosing operators.

Several pairs of operators can be used as enclosing operators. Any number of ‘`|`’ (vertical line) can be used as both infix operators and enclosing operators.

Some operators are always postfix: a ‘`^`’ followed by any ordinary operator (with no intervening whitespace) is considered to be a superscripted postfix operator. For example, ‘`^*`’ and ‘`^+`’ and ‘`^?`’ are available for use as part of the syntax of extended regular expressions. As a very special case, ‘`^T`’ is also considered to be a superscripted postfix operator, typically used to signify matrix transposition.

Finally, there are special operators such as juxtaposition and operators on dimensions and units. Juxtaposition may be a function application or an infix operator in Fortress. When the left-hand-side expression is a function, juxtaposition performs function application; when the left-hand-side expression is a number, juxtaposition conventionally performs multiplication; when the left-hand-side expression is a string, juxtaposition conventionally performs string concatenation. Fortress provides several operators on dimensions and units as described in Chapter ??.

14.1 Operator Names

14.2 Operator Names

Victor: This should refer to the appropriate sections of Lexical Structure Operator names are either operator tokens or special operator characters, or a few reserved words.

To support a rich mathematical notation, Fortress allows most Unicode characters that are specified to be mathematical operators to be used as operators in Fortress expressions, as well as these characters:

!	@	#	\$	%	*	+	-	=		:	<	>	/	?	^	~
->	-->	=>	==>	<=	>=	=/=	!!									
<<	<<<	>>	>>>	<->	<-/-	-/->	<=>	===								

In addition, a token that is made up of a mixture of uppercase letters and underscores (but no digits), does not begin or end with an underscore, and contains at least two different letters is also considered to be an operator:

MAX MIN

The above operators are rendered as: `MAX` `MIN`. (See Section ?? and Appendix ?? for detailed descriptions of operator names in Fortress.)

14.3 Operator Precedence and Associativity

Fortress specifies that certain operators have higher precedence than certain other operators and certain operators are associative, so that one need not use parentheses in all cases where operators are mixed in an expression. (See Appendix ?? for a detailed description of operator precedence and associativity in Fortress.) However, Fortress does not follow the practice of other programming languages in simply assigning an integer to each operator and then saying that the precedence of any two operators can be compared by comparing their assigned integers. Instead, Fortress relies on defining traditional groups of operators based on their meaning and shape, and specifies specific precedence relationships between some of these groups. If there is no specific precedence relationship between two operators, then parentheses must be used. For example, Fortress does not accept the expression $a + b \cup c$; one must write either $(a + b) \cup c$ or $a + (b \cup c)$. (Whether or not the result then makes any sense depends on what definitions have been made for the $+$ and \cup operators—see Chapter ??.)

Here are the basic principles of operator precedence and associativity in Fortress:

- Member selection (`.`) and method invocation (`.name(...)`) are not operators. They have higher precedence than any operator listed below.
- Subscripting (`[]` and any kind of subscripting operators, which can be any kind of enclosing operators), superscripting (`^`), and postfix operators have higher precedence than any operator listed below; within this group, these operations are left-associative (performed left-to-right).
- *Tight juxtaposition*, that is, juxtaposition without intervening whitespace, has higher precedence than any operator listed below. The associativity of tight juxtaposition is type-dependent; see Section 14.9.
- Next, *tight fractions*, that is, the use of the operator `'/'` with no whitespace on either side, have higher precedence than any operator listed below. The tight-fraction operator has no precedence compared with itself, so it is not permitted to be used more than once in a tight fraction without use of parentheses.
- *Loose juxtaposition*, that is, juxtaposition with intervening whitespace, has higher precedence than any operator listed below. The associativity of loose juxtaposition is type-dependent and is different from that for tight juxtaposition; see Section 14.9. Note that *lopsided juxtaposition* (having whitespace on one side but not the other) is a static error as described in Section 14.4.
- Prefix operators have higher precedence than any operator listed below. However, it is a static error for an operand of a loose prefix operator to be an operand of a tight infix operator.
- The infix operators are partitioned into certain traditional groups, as explained below. They have higher precedence than any operator listed below.

use of infix operator a has higher or equal precedence than a use of infix operator b , but that particular use of a is loose and that particular use of b is tight. Thus, for example, the expression $\sin x + y$ is permitted, but $\sin x+y$ is not permitted. Similarly, the expression $a \cdot b + c$ is permitted, as are $a \cdot b + c$ and $a \cdot b+c$, but $a \cdot b+c$ is not permitted. (The rule detects only the presence or absence of whitespace, not the amount of whitespace, so $a \cdot b + c$ is permitted. You have to draw the line somewhere.)

When in doubt, just use parentheses. If there's a problem, the compiler will (probably) let you know.

14.4 Operator Fixity

Multifix operators are not yet supported.

Most operators in Fortress can be used variously as prefix, postfix, infix, nofix, or multifix operators. (See Section 14.5 for a discussion of how infix operators may be chained or treated as multifix operators.)

Victor: I think that we shouldn't list "multifix" separately here. (Note that multifix does not appear in any of the tables.) Rather multifix is a way to interpret infix operators.

Some operators can be used in pairs as enclosing (bracketing) operators—see Section 14.6. The Fortress language dictates only the rules of syntax; whether an operator has a meaning when used in a particular way depends only on whether there is a definition in the program for that operator when used in that particular way (see Chapter ??).

The fixity of a non-enclosing operator is determined by context. To the left of such an operator we may find (1) a *primary tail* (described below), (2) another operator, or (3) a comma, semicolon, or left encloser. To the right we may find (1) a *primary front* (described below), (2) another operator, (3) a comma, semicolon, or right encloser, or (4) a line break. A primary tail is an identifier, a literal, a right encloser, or a superscripted postfix operator (exponent operator). A primary front is an identifier, a literal, or a left encloser.

A primary expression is an identifier, a literal, an expression enclosed by matching enclosers, a field selection, or an expression followed by a postfix operator.

Considered in all combinations, this makes twelve possibilities. In some cases one must also consider whether or not whitespace separates the operator from what lies on either side. The rules of operator fixity are specified by Figure 14.1, where the center column indicates the fixity that results from the left and right context specified by the other columns.

A case described in the center column of the table as an **error** is a static error; for such cases, the fixity mentioned in parentheses is the recommended treatment of the operator for the purpose of attempting to continuing the parse in search of other errors.

The table may seem complicated, but it all boils down to a couple of practical rules of thumb:

1. Any operator can be prefix, postfix, infix, or nofix.
2. An infix operator can be *loose* (having whitespace on both sides) or *tight* (having whitespace on neither side), but it mustn't be *lopsided* (having whitespace on one side but not the other).
3. A postfix operator must have no whitespace before it and must be followed (possibly after some whitespace) by a comma, semicolon, right encloser, or line break.

14.5 Chained and Multifix Operators

Certain infix mathematical operators that are traditionally regarded as *relational* operators, delivering boolean results, may be *chained*. For example, an expression such as $A \subseteq B \subset C \subseteq D$ is treated as being equivalent to $(A \subseteq B) \wedge (B \subset C) \wedge (C \subseteq D)$ except that the expressions B and C are evaluated only once (which matters only if they have side effects such as writes or input/output actions). Similarly, the expression $A \subseteq B = C \subset D$ is treated as being equivalent to $(A \subseteq B) \wedge (B = C) \wedge (C \subset D)$, except that B and C are evaluated only once. Fortress restricts such chaining to a mixture of equivalence operators and ordering operators; if a chain contains two or more ordering operators, then they must be of the same kind and have the same sense of monotonicity; for example, neither

left context	whitespace	operator fixity	whitespace	right context
primary tail	yes	infix	yes	primary front
	yes	error (infix)	no	
	no	postfix	yes	
	no	infix	no	
primary tail	yes	infix	yes	operator
	yes	error (infix)	no	
	no	postfix	yes	
	no	infix	no	
primary tail	yes	error (postfix)		, ; right enclosure
	no	postfix		
primary tail	yes	infix		line break
	no	postfix		
operator		prefix		primary front
operator		prefix		operator
operator		error (nofix)		, ; right enclosure
operator		error (nofix)		line break
, ; left enclosure		prefix		primary front
, ; left enclosure		prefix		operator
, ; left enclosure		nofix		, ; right enclosure
, ; left enclosure		error (prefix)		line break

Figure 14.1: Operator Fixity (I)

$A \subseteq B \leq C$ nor $A \subseteq B \supset C$ is permitted. This transformation is done before type checking. In particular, it is done even if these operators do not return boolean values, and the resulting expression is checked for type correctness. (See Section ?? for a detailed description of which operators may be chained.)

Any infix operator that does not chain may be treated as *multifix*. If $n - 1$ occurrences of the same operator separate n operands where $n \geq 3$, then the compiler first checks to see whether there is a definition for that operator that will accept n arguments. If so, that definition is used; if not, then the operator is treated as left-associative and the compiler looks for a two-argument definition for the operator to use for each occurrence. As an example, the cartesian product $S_1 \times S_2 \times \cdots \times S_n$ of n sets may usefully be defined as a multifix operator, but ordinary addition $p + q + r + s$ is normally treated as $((p + q) + r) + s$.

14.6 Enclosing Operators

These operators are always used in pairs as enclosing operators:

	(/ /)	(\ \)	
[]	[/ /]		[* *]
{ }	{/ /}	{\ \}	{ * *}
	</ />	<\ \>	
	<</ />>	<<\ \>>	

(ASCII encodings are shown here; they all correspond to particular single Unicode characters.) There are other pairs as well, such as `[]` and `[]` and multicharacter enclosing operators described in Section ?? . Note that the pairs `()` and `{\ \}` (also known as `[]`) are not operators; they play special roles in the syntax of Fortress, and their behavior cannot be redefined by a library. The bracket pairs that may be used as enclosing operators are described in Section ?? .

Any number of ‘|’ (vertical line) may also be used in pairs as enclosing operators but there is a trick to it, because on the face of it you can’t tell whether any given occurrence is a left enclosure or a right enclosure. Again, context is

left context	whitespace	operator fixity	whitespace	right context
primary tail	yes	infix	yes	primary front
	yes	left encloser	no	
	no	right encloser	yes	
	no	infix	no	
primary tail	yes	infix	yes	operator
	yes	left encloser	no	
	no	right encloser	yes	
	no	infix	no	
primary tail	yes	error (right encloser)		, ; right encloser
	no	right encloser		
primary tail	yes	infix		line break
	no	right encloser		
operator		error (left encloser)	yes	primary front
		left encloser	no	
operator		error (left encloser)	yes	operator
		left encloser	no	
operator		error (nofix)		, ; right encloser
operator		error (nofix)		line break
, ; left encloser		left encloser		primary front
, ; left encloser		left encloser		operator
, ; left encloser		nofix		, ; right encloser
, ; left encloser		error (left encloser)		line break

Figure 14.2: Operator Fixity (II)

used to decide, this time according to Figure 14.2.

This is very similar to Figure 14.1 in Section 14.4; a rough rule of thumb is that if an ordinary operator would be considered a prefix operator, then one of these will be considered a left encloser; and if an ordinary operator would be considered a postfix operator, then one of these will be considered a right encloser.

In this manner, one may use $|\dots|$ for absolute values and $||\dots||$ for matrix norms.

14.7 Conditional Operators

If a binary operator other than ‘:’ and ‘::’ is immediately followed by a ‘:’ then it is *conditional*: evaluation of the right-hand operand cannot begin until evaluation of the left-hand operand has completed, and whether or not the right-hand operand is evaluated may depend on the value of the left-hand operand. If the left-hand operand throws an exception, then the right-hand operand is not evaluated.

The Fortress standard libraries define several conditional operators on boolean values including \wedge : and \vee :

See Section ?? for a discussion of how conditional operators are declared.

14.8 Big Operators

A big operator application is either a *reduction expression* described in Section ?? or a *comprehension* described in Section 12.17.

The Fortress standard libraries define several big operators including \sum , \prod , and $\text{BIG } \wedge$.

See Section ?? for a discussion of how big operators are declared.

14.9 Juxtaposition

Qualified names are not yet supported.
--

Juxtaposition in Fortress may be a function call or a special infix operator. The Fortress standard libraries include several declarations of a `juxtaposition` operator.

When two expressions are juxtaposed, the juxtaposition is interpreted as follows: if the left-hand-side expression is a function, juxtaposition performs function application; otherwise, juxtaposition performs the `juxtaposition` operator application.

The manner in which a juxtaposition of three or more items must be associated requires type information and awareness of whitespace. (This is an inherent property of customary mathematical notation, which Fortress is designed to emulate where feasible.) Therefore a Fortress compiler must produce a provisional parse in which such multi-element juxtapositions are held in abeyance, then perform a type analysis on each element and use that information to rewrite the n -ary juxtaposition into a tree of binary juxtapositions.

All we need to know is whether each element of a juxtaposition has an arrow type. There are actually three legitimate possibilities for each element of a juxtaposition: (a) it has an arrow type, in which case it is considered to be a function element; (b) it has a type that is not an arrow type, in which case it is considered to be a non-function element. (c) it is an identifier that has no visible declaration, in which case it is considered to be a function element (and everything will work out okay if it turns out to be the name of an appropriate functional method).

The rules below are designed to forbid certain forms of notational ambiguity that can arise if the name of a functional method happens to be used also as the name of a variable. For example, suppose that trait T has a functional method of one parameter named n ; then in the code

```
do
  a: T = t
  n: ℤ = 14
  z = n a
end
```

it might not be clear whether the intended meaning was to invoke the functional method n on a or to multiply a by 14. The rules specify that such a situation is a static error.

The rules for reassociating a loose juxtaposition are as follows:

- First the loose juxtaposition is broken into nonempty chunks; wherever there is a non-function element followed by a function element, the latter begins a new chunk. Thus a chunk consists of some number (possibly zero) of functions followed by some number (possibly zero) of non-functions.
- It is a static error if any non-function element in a chunk is an unparenthesized identifier f and is followed by another non-function element whose type is such that f can be applied to that latter element as a functional method.
- The non-functions in each chunk, if any, are replaced by a single element consisting of the non-functions grouped left-associatively into binary juxtapositions.
- What remains in each chunk is then grouped right-associatively.
(Notice that, up to this point, no analysis of the types of newly constructed chunks is needed during this process.)
- It is a static error if an element of the original juxtaposition was the last element in its chunk before reassociation, the chunk was not the last chunk (and therefore the element in question is a non-function element), the element was an unparenthesized identifier f , and the type of the following chunk after reassociation is such that f can be applied to that following chunk as a functional method.
- If any element that remains has type `String`, then it is a static error if there is any pair of adjacent elements within the juxtaposition such that neither element is of type `String`.

- What remains is considered to be a binary or multifix application of the juxtaposition operator; if more than two elements remain, the application is handled in the same way as for any other multifix operator (see Section 14.5).

Here is an example: $n(n+1) \sin 3nx \log \log x$. Assuming that \sin and \log name functions in the usual manner and that n , $(n+1)$, and x are not functions, this loose juxtaposition splits into three chunks: $n(n+1)$, $\sin 3nx$, and $\log \log x$. The first chunk has only two elements and needs no further reassociation. In the second chunk, the non-functions $3nx$ are replaced by $((3n)x)$. In the third chunk, there is only one non-function, so that remains unchanged; the chunk is the right-associated to form $(\log(\log x))$. Assuming that no multifix definition of juxtaposition has been provided, the three chunks are left-associated, to produce the final interpretation $((n(n+1))(\sin((3n)x))) (\log(\log x))$. Now the original juxtaposition has been reduced to binary juxtaposition expressions.

A tight juxtaposition is always left-associated if it contains any dot (i.e., “.”) not within parentheses or some pair of enclosing operators. If such a tight juxtaposition begins with an identifier immediately followed by a dot then the maximal prefix of identifiers separated by dots (whitespace may follow but not precede the dots) is collected into a “dotted id chain”, which is subsequently partitioned into the longest prefix, if any, that is a qualified name, and the remaining the dots and identifiers, which are interpreted as selectors. If the last identifier in the dotted id chain is part of a selector (i.e., the entire dotted id chain is not a qualified name), and it is immediately followed by a left parenthesis, then the last selector together with the subsequent parenthesis-delimited expression is a method invocation.

A tight juxtaposition without any dots might not be entirely left-associated. Rather, it is considered as a nonempty sequence of elements: the front expression, any “math items”, and any postfix operator, and subject to reassociation as described below. A math item may be a subscripting, an exponentiation, or one of a few kinds of expressions. It is a static error if an exponentiation is immediately followed by a subscripting or an exponentiation.

The procedure for reassociation is as follows:

- For each expression element (i.e., not a subscripting, exponentiation or postfix operator), determine whether it is a function.
- If some function element is immediately followed by an expression element then, find the first such function element, and call the next element the argument. It is a static error if either the argument is not parenthesized, or the argument is immediately followed by a non-expression element. Otherwise, replace the function and argument with a single element that is the application of the function to the argument. This new element is an expression. Reassociate the resulting sequence (which is one element shorter).
- If there is any non-expression element (it cannot be the first element) then replace the first such element and the element immediately preceding it (which must be an expression) with a single element that does the appropriate operator application. This new element is an expression. Reassociate the resulting sequence (which is one element shorter).
- Otherwise, the sequence necessarily has only expression elements, only the last of which may be a function. If any element that remains has type String, then it is a static error if there is any pair of adjacent elements within the juxtaposition such that neither element is of type String.
- What remains is considered to be a binary or multifix application of the juxtaposition operator; if more than two elements remain, the application is handled in the same way as for any other multifix operator (see Section 14.5).

(Note that this process requires type analysis of newly created chunks all along the way.)

Here is an (admittedly contrived) example: $reduce(f)(a)(x+1)atan(x+2)$. Suppose that $reduce$ is a curried function that accepts a function f and returns a function that can be applied to an array a (the idea is to use the function f , which ought to take two arguments, to combine the elements of the array to produce an accumulated result).

The leftmost function is $reduce$, and the following element (f) is parenthesized, so the two elements are replaced with one: $(reduce(f))(a)(x+1)atan(x+2)$. Now type analysis determines that the element $(reduce(f))$ is a function.

The leftmost function is $(reduce(f))$, and the following element (a) is parenthesized, so the two elements are replaced with one: $((reduce(f))(a))(x+1)atan(x+2)$. Now type analysis determines that the element $((reduce(f))(a))$ is not a function.

The leftmost function is *atan*, and the following element $(x + 2)$ is parenthesized, so the two elements are replaced with one: $((\text{reduce}(f))(a))(x + 1)(\text{atan}(x + 2))$. Now type analysis determines that the element $(\text{atan}(x + 2))$ is not a function.

There are no functions remaining in the juxtaposition. Assuming that no multifix definition of juxtaposition has been provided, the remaining elements are left-associated:

$$(((\text{reduce}(f))(a))(x + 1))(\text{atan}(x + 2))$$

Now the original juxtaposition has been reduced to binary juxtaposition expressions.

14.10 Operator Declarations

Multifix operators and keyword parameters are not yet supported.

An operator declaration may appear anywhere a top-level function or method declaration may appear. Operator declarations are like other function or method declarations in all respects except that an operator declaration has *opr* and has an operator name (see Section 14.2 for a discussion of valid operator names) instead of an identifier. The precise placement of the operator name within the declaration depends on the fixity of the operator. Like other functionals, operators may have overloaded declarations (see Chapter ?? for a discussion of overloading). These overloads may be of the same or differing fixities.

Syntax:

<i>FnDecl</i>	::=	<i>FnMods?</i> <i>FnHeaderFront</i> <i>FnHeaderClause</i> (= <i>Expr</i>)?
<i>FnHeaderFront</i>	::=	<i>OpHeaderFront</i>
<i>OpHeaderFront</i>	::=	<i>opr</i> <i>BIG?</i> (<i>{</i> \mapsto <i>LeftEncloser</i> <i>Encloser</i>) <i>StaticParams?</i> <i>Params?</i> (<i>RightEncloser</i> <i>Encloser</i>) <i>opr</i> <i>ValParam</i> (<i>Op</i> <i>ExponentOp</i>) <i>StaticParams?</i> <i>opr</i> <i>BIG?</i> (<i>Op</i> \wedge <i>Encloser</i> \sum \prod) <i>StaticParams?</i> <i>ValParam</i>
<i>MdDecl</i>	::=	<i>MdDef</i> <i>abstract?</i> <i>MdMods?</i> <i>MdHeaderFront</i> <i>FnHeaderClause</i>
<i>MdHeaderFront</i>	::=	<i>OpMdHeaderFront</i>
<i>OpMdHeaderFront</i>	::=	<i>opr</i> <i>BIG?</i> (<i>{</i> \mapsto <i>LeftEncloser</i> <i>Encloser</i>) <i>StaticParams?</i> <i>Params?</i> (<i>RightEncloser</i> <i>Encloser</i>) (<i>:=</i> (<i>SubscriptAssignParam</i>))? <i>opr</i> <i>ValParam</i> (<i>Op</i> <i>ExponentOp</i>) <i>StaticParams?</i> <i>opr</i> <i>BIG?</i> (<i>Op</i> \wedge <i>Encloser</i> \sum \prod) <i>StaticParams?</i> <i>ValParam</i>
<i>SubscriptAssignParam</i>	::=	<i>Varargs</i> <i>Param</i>

An operator declaration has one of eight forms: multifix operator declaration, infix operator declaration, prefix operator declaration, postfix operator declaration, nofix operator declaration, bracketing operator declaration, subscripting operator method declaration, and subscripted assignment operator method declaration. Each is invoked according to specific rules of syntax. An operator method declaration must be a functional method declaration, a subscripting operator method declaration, or a subscripted assignment operator method declaration.

14.10.1 Multifix Operator Declarations

A multifix operator declaration has *opr* and then an operator name where a functional declaration would have an identifier. The declaration must not have any keyword parameters, and must be capable of accepting at least two arguments. It is permissible to use a *varargs* parameter; in fact, this is a good way to define a multifix operator. Static parameters (described in Chapter ??) may also be present, between the operator and the parameter list.

An expression consisting of a multifix operator applied to an expression will invoke a multifix operator declaration. The compiler considers all multifix operator declarations for that operator that are both accessible and applicable, and

the most specific operator declaration is chosen according to the usual rules for overloaded functionals. The invocation will pass more than two arguments.

A multifix operator declaration may also be invoked by an infix, a prefix, or nofix (but not a postfix) operator application if the declaration is applicable.

Example:

Example is commented out because it is not supported nor run by the interpreter yet.

14.10.2 Infix Operator Declarations

An infix operator declaration has `opr` and then an operator name where a functional declaration would have an identifier. The declaration must have two value parameter, which must not be a keyword parameter or varargs parameter. Static parameters may also be present, between the operator and the parameter list.

An expression consisting of an infix operator applied to an expression will invoke an infix operator declaration. The compiler considers all infix and multifix operator declarations for that operator that are both accessible and applicable, and the most specific operator declaration is chosen according to the usual rules for overloaded functionals.

Note that superscripting (\wedge) may be defined using an infix operator declaration even though it has very high precedence and cannot be used as a multifix operator.

Example:

14.10.3 Prefix Operator Declarations

A prefix operator declaration has `opr` and then an operator name where a functional declaration would have an identifier. The declaration must have one value parameter, which must not be a keyword parameter or varargs parameter. Static parameters may also be present, between the operator and the parameter list.

An expression consisting of a prefix operator applied to an expression will invoke a prefix operator declaration. The compiler considers all prefix and multifix operator declarations for that operator that are both accessible and applicable, and the most specific operator declaration is chosen according to the usual rules for overloaded functionals.

Example:

14.10.4 Postfix Operator Declarations

A postfix operator declaration has `opr` where a functional declaration would have an identifier; the operator name itself *follows* the parameter list. The declaration must have one value parameter, which must not be a keyword parameter or varargs parameter. Static parameters may also be present, between `opr` and the parameter list.

An expression consisting of a postfix operator applied to an expression will invoke a postfix operator declaration. The compiler considers all postfix operator declarations for that operator that are both accessible and applicable, and the most specific operator declaration is chosen according to the usual rules for overloaded functionals.

Example:

14.10.5 Nofix Operator Declarations

A nofix operator declaration has `opr` and then an operator name where a functional declaration would have an identifier. The declaration must have no parameters.

An expression consisting only of a nofix operator will invoke a nofix operator declaration. The compiler considers all nofix and multifix operator declarations for that operator that are both accessible and applicable, and the most specific operator declaration is chosen according to the usual rules for overloaded functionals.

Uses for nofix operators are rare, but those rare examples are very useful. For example, if the `@` operator is used to construct subscripting ranges, and it is the nofix declaration of `@` that allows a lone `@` to be used as a subscript:

14.10.6 Bracketing Operator Declarations

A bracketing operator declaration has `opr` where a functional declaration would have an identifier. The value parameter list, rather than being surrounded by parentheses, is surrounded by the brackets being defined. A bracketing operator declaration may have any number of parameters, keyword parameters, and `varargs` parameters in the value parameter list. Static parameters may also be present, between `opr` and the parameter list. Any paired Unicode brackets may be so defined *except* ordinary parentheses and white square brackets.

An expression consisting of zero or more comma-separated expressions surrounded by a bracket pair will invoke a bracketing operator declaration. The compiler considers all bracketing operator declarations for that type of bracket pair that are both accessible and applicable, and the most specific operator declaration is chosen according to the usual rules for overloaded functionals. For example, the expression $\langle p, q \rangle$ might invoke the following bracketing method declaration:

14.11 Overview of Operators in the Fortress Standard Libraries

The operators in this section are not tested nor run by the interpreter yet.

This section provides a high-level overview of the operators in the Fortress standard libraries. See Appendix ?? for the detailed rules for the operators provided by the Fortress standard libraries.

14.11.1 Prefix Operators

For all standard numeric types, the prefix operator `+` simply returns its argument and the prefix operator `-` returns the negative of its argument.

The operator \neg is the logical NOT operator on boolean values and boolean intervals.

The operator \neg computes the bitwise NOT of an integer.

Big operators such as \sum begin a *reduction expression* (Section ??). The big operators include \sum (summation) and \prod (product), along with \bigcap , \bigcup , \bigwedge , \bigvee , \bigvee , \bigoplus , \bigotimes , \boxplus , \boxtimes , \boxtimes , MAX , MIN , and so on.

14.11.2 Postfix Operators

The operator `!` computes factorial; the operator `!!` computes double factorials. They may be applied to a value of any integral type and produces a result of the same type.

When applied to a floating-point value x , $x!$ computes $\Gamma(1 + x)$, where Γ is the Euler gamma function.

14.11.3 Enclosing Operators

When used as left and right enclosing operators, `| |` computes the absolute value or magnitude of any number, and is also used to compute the number of elements in an aggregate, for example the cardinality of a set or the length of a list. Similarly, `||` is used to compute the norm of a vector or matrix.

The floor operator $\lfloor \cdot \rfloor$ and ceiling operator $\lceil \cdot \rceil$ may be applied to any standard integer, rational, or real value (their behavior is trivial when applied to integers, of course). The operators hyperfloor $\llbracket x \rrbracket = 2^{\lfloor \log_2 x \rfloor}$, hyperceiling $\lceil\lceil x \rceil\rceil = 2^{\lceil \log_2 x \rceil}$, hyperhyperfloor $\llbracket\lceil x \rceil\rrbracket = 2^{\lfloor \log_2 x \rfloor}$, and hyperhyperceiling $\lceil\lfloor x \rfloor\rceil = 2^{\lceil \log_2 x \rceil}$ are also available.

14.11.4 Exponentiation

Given two expressions e and e' denoting numeric quantities v and v' that are not vectors or matrices, the expression $e^{e'}$ denotes the quantity obtained by raising v to the power v' . This operation is defined in the usual way on numerals.

Given an expression e denoting a vector and an expression e' denoting a value of type \mathbb{Z} , the expression $e^{e'}$ denotes repeated vector multiplication of e by itself e' times.

Given an expression e denoting a square matrix and an expression e' denoting a value of type \mathbb{Z} , the expression $e^{e'}$ denotes repeated matrix multiplication of e by itself e' times.

Eric: Do we also want to support the matrix exponential via this operator? Doing so requires having a special constant value for e .

14.11.5 Superscript Operators

The superscript operator T transposes a matrix. It also converts a column vector to a row vector or a row vector to a column vector.

14.11.6 Subscript Operators

Subscripting of arrays and other aggregates is written using square brackets:

<code>a[i]</code>	<i>is displayed as</i>	a_i	i th element of one-dimensional array a
<code>m[i, j]</code>	<i>is displayed as</i>	m_{ij}	i, j th element of two-dimensional matrix m
<code>space[i, j, k]</code>	<i>is displayed as</i>	$space_{ijk}$	i, j, k th element of three-dimensional array $space$
<code>a[3] := 4</code>	<i>is displayed as</i>	$a_3 := 4$	assign 4 to the third element of mutable array a
<code>m["foo"]</code>	<i>is displayed as</i>	$m_{\text{“foo”}}$	fetch the entry associated with string “foo” from map m

14.11.7 Multiplication, Division, Modulo, and Remainder Operators

For most integer, rational, floating-point, complex, and interval expressions, multiplication can be expressed using any of \cdot or \times or simply juxtaposition. However, the \times operator is used to express the shape of matrices, so if expressions using multiplication are used in expressing the shape of a matrix, it may be necessary to avoid the use of \times to express multiplication, or to use parentheses.

For integer, rational, floating-point, complex, and interval expressions, division is expressed by $/$. When the operator $/$ is used to divide one integer by another, the result is rational. The operator \div performs truncating integer division: $m \div n = \text{signum}(\frac{m}{n}) \lfloor \lceil \frac{m}{n} \rceil \rfloor$. The operator `REM` gives the remainder from such a truncating division: $m \text{ REM } n = m - n(m \div n)$. The operator `MOD` gives the remainder from a floor division: $m \text{ MOD } n = m - n \lfloor \frac{m}{n} \rfloor$; when $n > 0$ this is the usual modulus computation that evaluates integer m to an integer k such that $0 \leq k < n$ and n evenly divides $m - k$.

The special operators `DIVREM` and `DIVMOD` each return a pair of values, the quotient and the remainder; $m \text{ DIVREM } n$ returns $(m \div n, m \text{ REM } n)$ while $m \text{ DIVMOD } n$ returns $(\lfloor \frac{m}{n} \rfloor, m \text{ MOD } n)$.

Multiplication of a vector or matrix by a scalar is done with juxtaposition, as is multiplication of a vector by a matrix (on either side). Vector dot product is expressed by \cdot and vector cross product by \times . Division of a matrix or vector by a scalar may be expressed using $/$.

The syntactic interaction of juxtaposition, \cdot , \times , and $/$ is subtle. See Section 14.3 for a discussion of the relative precedence of these operations and how precedence may depend on the use of whitespace.

The handling of overflow depends on the type of the number produced. For integer results, overflow throws an `IntegerOverflow`. Rational computations do not overflow. For floating-point results, overflow produces $+\infty$ or $-\infty$ according to the rules of IEEE 754. For intervals, overflow produces an appropriate containing interval.

Underflow is relevant only to floating-point computations and is handled according to the rules of IEEE 754.

The handling of division by zero depends on the type of the number produced. For integer results, division by zero throws a `DivisionByZero`. For rational results, division by zero produces $1/0$. For floating-point results, division by zero produces a NaN value according to the rules of IEEE 754. For intervals, division by zero produces an appropriate containing interval (which under many circumstances will be the interval of all possible real values and infinities).

Wraparound multiplication on fixed-size integers is expressed by $\dot{\times}$. Saturating multiplication on fixed-size integers is expressed by \boxtimes or \boxtimes . These operations do not overflow.

Ordinary multiplication and division of floating-point numbers always use the IEEE 754 “round to nearest” rounding mode. This rounding mode may be emphasized by using the operators \otimes (or \odot) and \oslash . Multiplication and division in “round toward zero” mode may be expressed with \boxtimes (or \boxdiv) and \boxdiv . Multiplication and division in “round toward positive infinity” mode may be expressed with \triangleup (or \triangleup) and \triangleup . Multiplication and division in “round toward negative infinity” mode may be expressed with ∇ (or ∇) and ∇ .

14.11.8 Addition and Subtraction Operators

Addition and subtraction are expressed with $+$ and $-$ on all numeric quantities, including intervals, as well as vectors and matrices.

The handling of overflow depends on the type of the number produced. For integer results, overflow throws an `IntegerOverflow`. Rational computations do not overflow. For floating-point results, overflow produces $+\infty$ or $-\infty$ according to the rules of IEEE 754. For intervals, overflow produces an appropriate containing interval.

Underflow is relevant only to floating-point computations and is handled according to the rules of IEEE 754.

Wraparound addition and subtraction on fixed-size integers are expressed by $\dot{+}$ and $\dot{-}$. Saturating addition and subtraction on fixed-size integers are expressed by \boxplus and \boxminus . These operations do not overflow.

Ordinary addition and subtraction of floating-point numbers always use the IEEE 754 “round to nearest” rounding mode. This rounding mode may be emphasized by using the operators \oplus and \ominus . Addition and subtraction in “round toward zero” mode may be expressed with \boxplus and \boxminus . Addition and subtraction in “round toward positive infinity” mode may be expressed with \triangleup and \triangleleft . Addition and subtraction in “round toward negative infinity” mode may be expressed with ∇ and ∇ .

The construction $x \pm y$ produces the interval $\langle x - y, x + y \rangle$.

14.11.9 Intersection, Union, and Set Difference Operators

Sets support the operations of intersection \cap , union \cup , disjoint union \uplus , set difference \setminus , and symmetric set difference \ominus . Disjoint union throws `DisjointUnionError` if the arguments are in fact not disjoint.

Intervals support the operations of intersection \cap , union \cup , and interior hull \sqcup . The operation \cap returns a pair of intervals; if the intersection of the arguments is a single contiguous span of real numbers, then the first result is an interval representing that span and the second result is an empty interval, but if the intersection is two spans, then two (disjoint) intervals are returned. The operation \cup returns a pair of intervals; if the arguments overlap, then the first result is the union of the two intervals and the second result is an empty interval, but if the arguments are disjoint, they are simply returned as is.

14.11.10 Minimum and Maximum Operators

The operator `MAX` returns the larger of its two operands, and `MIN` returns the smaller of its two operands.

For floating-point numbers, if either argument is a NaN then NaN is returned. The floating-point operations `MAXNUM` and `MINNUM` behave similarly except that if one argument is NaN and the other is a number, the number is returned. For all four of these operators, when applied to floating-point values, -0 is considered to be smaller than $+0$.

14.11.11 GCD, LCM, and CHOOSE Operators

The infix operator `GCD` computes the greatest common divisor of its integer operands, and `LCM` computes the least common multiple. The operator `CHOOSE` computes binomial coefficients: $n \text{ CHOOSE } k = \binom{n}{k} = \frac{n!}{k!(n-k)!}$.

The expression $e \neq e'$ is semantically equivalent to the expression $\neg(e = e')$.

14.11.12 Comparisons Operators

Unless otherwise noted, the operators described in this section produce boolean (*true/false*) results.

The operators $<$, \leq , \geq , and $>$ are used for numerical comparisons and are supported by integer, rational, and floating-point types. Comparison of rational values throws `RationalComparisonError` if either argument is the rational infinity $1/0$ or the rational indefinite $0/0$. Comparison of floating-point values throws `FloatingComparisonError` if either argument is a NaN.

The operators $<$, \leq , \geq , and $>$ may also be used to compare characters (according to the numerical value of their Unicode codepoint values) and strings (lexicographic order based on the character ordering). They also use lexicographic order when used to compare lists whose elements support these same comparison operators.

When $<$, \leq , \geq , and $>$ are used to compare numerical intervals, the result is a boolean interval. The functions *possibly* and *certainly* are useful for converting boolean intervals to boolean values for testing. Thus *possibly*($x > y$) is true if and only if there is some value in the interval x that is greater than some value in the interval y , while *certainly*($x > y$) is true if and only if x and y are nonempty and every value in x is greater than every value in y .

The operators \subset , \subseteq , \supseteq , and \supset may be used to compare sets or intervals regarded as sets.

The operator \in may be used to test whether a value is a member of a set, list, array, interval, or range.

14.11.13 Logical Operators

The following binary operators may be used on boolean values:

\wedge	AND
\vee	inclusive OR
$\underline{\vee}$ or \oplus or \neq	exclusive OR
\equiv or \leftrightarrow or $=$	equivalence (if and only if)
\rightarrow	IMPLIES
$\overline{\wedge}$	NAND
$\overline{\vee}$	NOR

These same operators may also be applied to boolean intervals to produce boolean interval results.

The following operators may be used on integers to perform “bitwise” operations:

$\&$	bitwise AND
$\&\&$	bitwise inclusive OR
$\&\&\&$	bitwise exclusive OR

The prefix operator \neg computes the bitwise NOT of an integer.

14.11.14 Conditional Operators

If $p(x)$ and $q(x)$ are expressions that produce boolean results, the expression $p(x) \wedge q(x)$ computes the logical AND of those two results by first evaluating $p(x)$. If the result of $p(x)$ is *true*, then $q(x)$ is also evaluated, and its result becomes the result of the entire expression; but if the result of $p(x)$ is *false*, then $q(x)$ is not evaluated, and the result of the entire expression is *false* without further ado. (Similarly, evaluating the expression $p(x) \vee q(x)$ does not evaluate $q(x)$ if the result of $p(x)$ is *true*.) Contrast this with the expression $p(x) \wedge q(x)$ (with no colon), which evaluates both $p(x)$ and $q(x)$, in no specified order and possibly in parallel.

Chapter 15

Overloading and Dispatch

Keyword and varargs parameters are not yet supported.

Fortress allows functions and methods (collectively called *procedures*) to be *overloaded*. That is, there may be multiple declarations for the same procedure name visible in a single scope (which may include inherited method declarations), and several of them may be applicable to any particular procedure call. This raises the question of which definition will actually be executed for any given procedure invocation. The simple answer is that the *dynamically most specific applicable visible* definition is chosen. That is, one first considers the set of all definitions that are *dynamically visible*; then, among those in the set that are *applicable* to the given argument values, the *most specific* one is chosen.

At compile time, typechecking performs a related series of tests: the type of a procedure call is the return type of the *statically most specific applicable visible* declaration. That is, one first considers the set of all definitions that are *statically visible*; then, among those in the set that are *applicable* to the types of the given arguments, the *most specific* one is chosen.

```
trait T extends {A, B} end
```

In this chapter, we describe how to determine which declarations are *visible* to a particular procedure call (both statically and dynamically); how to further determine which of these are *applicable* to that procedure call (both statically and dynamically); and how the most specific one is chosen. We also introduce some rules for writing procedure declarations that ensure the soundness of the type system, as well as the existence and uniqueness of a most specific applicable declaration (at compile time) or definition (at run time). The result is that if a program successfully compiles, then procedure calls always succeed and are never ambiguous.

Section 15.1 introduces some terminology and notation. In Section 15.3, we show how to determine which declarations are applicable to a *named procedure call* (a function call described in Section 12.4 or a naked method invocation described in Section ??) when all declarations have only ordinary parameters (without varargs or keyword parameters). We discuss how to handle dotted method calls (described in Section ??) in Section 15.4, and declarations with varargs and keyword parameters in Section 15.5. Determining which declaration is applied, if several are applicable, is discussed in Section 15.6.

15.1 Principles of Overloading

Fortress allows multiple procedure declarations of the same name to be declared in a single scope. However, recall from Chapter ?? the following shadowing rules:

- dotted method declarations shadow top-level function declarations with the same name, and
- dotted method declarations provided by a trait or object declaration or object expression shadow functional method declarations with the same name that are provided by a different trait or object declaration or object expression.

Also, note that a trait or object declaration or object expression must not have a functional method declaration and a dotted method declaration with the same name, either directly or by inheritance. Therefore, top-level functions can overload with other top-level functions and functional methods, dotted methods with other dotted methods, and functional methods with other functional methods and top-level functions. It is a static error if any top-level function declaration is more specific than any functional method declaration. If a top-level function declaration is overloaded with a functional method declaration, the top-level function declaration must not be more specific than the functional method declaration.

Operator declarations with the same name but different fixity are not a valid overloading; they are unambiguous declarations. An operator method declaration whose name is one of the operator parameters (described in Section ??) of its enclosing trait or object may be overloaded with other operator declarations in the same component. Therefore, such an operator method declaration must satisfy the overloading rules (described in Chapter ??) with every operator declaration in the same component.

This restriction will be relaxed.

Recall from Chapter 5 that we write $T \preceq U$ when T is a subtype of U , and $T \prec U$ when $T \preceq U$ and $U \not\preceq T$.

15.2 Visibility to Named Procedure Calls

15.3 Applicability to Named Procedure Calls

In this section, we show how to determine which declarations are applicable to a named procedure call when all declarations have only ordinary parameters (i.e., neither varargs nor keyword parameters).

For the purpose of defining applicability, a named procedure call can be characterized by the name of the procedure and its argument type. Recall that a procedure has a single parameter, which may be a tuple (a dotted method has a receiver as well). We abuse notation by using *static call* $f(A)$ to refer to a named procedure call with name f and whose argument has static type A , and *dynamic call* $f(X)$ to refer to a named procedure call with We use $\text{call } f(C)$ to refer to a named procedure call with name f and whose argument, when evaluated, has dynamic type C . (Note that if the type system is sound—and we certainly hope that it is!—then $X \preceq A$ for all well-typed calls to f .) We use the term $\text{call } f(C)$ to refer to static and dynamic calls collectively. We assume throughout this chapter that all static variables in procedure calls have been instantiated or inferred.

We also use *function declaration* $f(P) : U$ to refer to a function declaration with function name f , parameter type P , and return type U .

For method declarations, we must take into account the self parameter, as follows:

A *dotted method declaration* $P_0.f(P) : U$ is a dotted method declaration with name f , where P_0 is the trait or object type in which the declaration appears, P is the parameter type, and U is the return type. (Note that despite the suggestive notation, a dotted method declaration does not explicitly list its self parameter.)

A *functional method declaration* $f(P) : U$ with *self parameter at i* is a functional method declaration with name f , with the parameter `self` in the i th position of the parameter type P , and return type U . Note that the static type of the self parameter is the trait or object trait type in which the declaration $f(P) : U$ occurs. In the following, we will use P_i to refer to the i th element of P .

We elide the return type of a declaration, writing $f(P)$ and $P_0.f(P)$, when the return type is not relevant to the discussion. Note that static parameters may appear in the types P_0 , P , and U .

A declaration $f(P)$ is *applicable* to a call $f(C)$ if the call is in the scope of the declaration and $C \preceq P$. (See Chapter ?? for the definition of scope.) If the parameter type P includes static parameters, they are inferred as described in Chapter 17 before checking the applicability of the declaration to the call.

Note that a named procedure call $f(C)$ may invoke a dotted method declaration if the declaration is provided by the trait or object enclosing the call. To account for this, let C_0 be the trait or object declaration immediately enclosing the call. Then we consider a named procedure call $f(C)$ as $C_0.f(C)$ if C_0 provides dotted method declarations applicable to $f(C)$, and use the rule for applicability to dotted method calls (described in Section 15.4) to determine which declarations are applicable to $C_0.f(C)$.

15.4 Applicability to Dotted Method Calls

Dotted method applications can be characterized similarly to named procedure applications, except that, analogously to dotted method declarations, we use C_0 to denote the dynamic type of the receiver object, and, as for named procedure calls, C to denote the dynamic type of the argument of a dotted method call. We write $C_0.f(C)$ to refer to the call.

A dotted method declaration $P_0.f(P)$ is *applicable* to a dotted method call $C_0.f(C)$ if $C_0 \preceq P_0$ and $C \preceq P$. If the types P_0 and P include static parameters, they are inferred as described in Chapter 17 before checking the applicability of the declaration to the call.

15.5 Applicability for Procedures with Varargs and Keyword Parameters

The basic idea for handling varargs and keyword parameters is that we can think of a procedure declaration that has such parameters as though it were (possibly infinitely) many declarations, one for each set of arguments it may be called with. In other words, we expand these declarations so that there exists a declaration for each number of arguments that can be passed to it.

A declaration with a varargs parameter corresponds to an infinite number of declarations, one for every number of arguments that may be passed to the varargs parameter. In practice, we can bound that number by the maximum number of arguments that the procedure is called with anywhere in the program (in other words, a given program will contain only a finite number of calls with different numbers of arguments). The expansion described here is a conceptual one to simplify the description of the semantics; we do not expect a real implementation to actually expand these declarations at compile time. For example, the following declaration:

$$f(x : \mathbb{Z}, y : \mathbb{Z}, z : \mathbb{Z} \dots) : \mathbb{Z}$$

would be expanded into:

$$\begin{aligned} &f(x : \mathbb{Z}, y : \mathbb{Z}) : \mathbb{Z} \\ &f(x : \mathbb{Z}, y : \mathbb{Z}, z_1 : \mathbb{Z}) : \mathbb{Z} \\ &f(x : \mathbb{Z}, y : \mathbb{Z}, z_1 : \mathbb{Z}, z_2 : \mathbb{Z}) : \mathbb{Z} \\ &f(x : \mathbb{Z}, y : \mathbb{Z}, z_1 : \mathbb{Z}, z_2 : \mathbb{Z}, z_3 : \mathbb{Z}) : \mathbb{Z} \\ &\dots \end{aligned}$$

A declaration with a varargs parameter is applicable to a call if any one of the expanded declarations is applicable.

15.6 Overloading Resolution

Victor: Does this paragraph, other than the last sentence, really belong in this section?

To evaluate a given procedure call, it is necessary to determine which procedure declaration to dispatch to. To do so, we consider the declarations that are applicable to that call at run time. If there is exactly one such declaration, then the call dispatches to that declaration. If there is no such declaration, then the call is *undefined*, which is a static error. (However, see Section ?? for how coercion may add to the set of applicable declarations.) If multiple declarations are applicable to the call at run time, then we choose an arbitrary declaration among the declarations such that no other applicable declaration is more specific than them.

We use the subtype relation to compare parameter types to determine a more specific declaration. Formally, a declaration $f(P)$ is *more specific* than a declaration $f(Q)$ if $P \prec Q$. Similarly, a declaration $P_0.f(P)$ is more specific than a declaration $Q_0.f(Q)$ if $P_0 \prec Q_0$ and $P \prec Q$. (See Section ?? for how coercion changes the definition of “more specific”.) Restrictions on the definition of overloaded procedures (see Chapter ??) guarantee that among all applicable declarations, one is more specific than all the others. If the declarations include static parameters, they are inferred as described in Chapter 17 before comparing their parameter types to determine which declaration is more specific.

Chapter 16

Coercion

Chapter 17

Type Inference

Chapter 18

Components and APIs

Appendix A

Simplified Grammar for Application Programmers and Library Writers

A.1 Components and API

File	::=	CompilationUnit Imports [?] Exports Decls [?] Imports [?] AbsDecls Imports AbsDecls [?]
CompilationUnit	::=	Component Api
Component	::=	component DottedId Imports [?] Exports Decls [?] end
Api	::=	api DottedId Imports [?] AbsDecls [?] end
Imports	::=	Import ⁺
Import	::=	import ImportFrom import AliasedDottedIds
ImportFrom	::=	* (except Names) [?] from DottedId AliasedNames from DottedId
Names	::=	Name { NameList }
NameList	::=	Name (, Name) [*]
AliasedNames	::=	AliasedName { AliasedNameList }
AliasedName	::=	Id (as DottedId) [?] opr Op (as Op) [?] opr LeftEncloser RightEncloser (as LeftEncloser RightEncloser) [?]
AliasedNameList	::=	AliasedName (, AliasedName) [*]
AliasedDottedIds	::=	AliasedDottedId { AliasedDottedIdList }
AliasedDottedId	::=	DottedId (as DottedId) [?]
AliasedDottedIdList	::=	AliasedDottedId (, AliasedDottedId) [*]
Exports	::=	Export ⁺
Export	::=	export DottedIds
DottedIds	::=	DottedId

		{ DottedIdList }
DottedIdList	::=	DottedId (, DottedId)*

A.2 Top-level Declarations

Decls	::=	Decl ⁺
Decl	::=	TraitDecl ObjectDecl VarDecl FnDecl
AbsDecls	::=	AbsDecl ⁺
AbsDecl	::=	AbsTraitDecl AbsObjectDecl AbsVarDecl AbsFnDecl

A.3 Trait Declaration

TraitDecl	::=	TraitHeader GoInATrait [?] end
TraitHeader	::=	TraitMods [?] trait Id StaticParams [?] Extends [?] TraitClauses [?]
TraitClauses	::=	TraitClause ⁺
TraitClause	::=	Excludes Comprises
GoInATrait	::=	GoFrontInATrait GoBackInATrait [?] GoBackInATrait
GoFrontInATrait	::=	GoesFrontInATrait ⁺
GoesFrontInATrait	::=	AbsFldDecl GetterSetterDecl
GoBackInATrait	::=	GoesBackInATrait ⁺
GoesBackInATrait	::=	MdDecl
AbsTraitDecl	::=	TraitHeader AbsGoInATrait [?] end
AbsGoInATrait	::=	AbsGoFrontInATrait AbsGoBackInATrait [?] AbsGoBackInATrait
AbsGoFrontInATrait	::=	AbsGoesFrontInATrait ⁺
AbsGoesFrontInATrait	::=	ApiFldDecl AbsGetterSetterDecl
AbsGoBackInATrait	::=	AbsGoesBackInATrait ⁺
AbsGoesBackInATrait	::=	AbsMdDecl AbsCoercion

A.4 Object declaration

ObjectDecl	::=	ObjectHeader GoInAnObject [?] end
ObjectHeader	::=	ObjectMods [?] object Id StaticParams [?] ObjectValParam [?] Extends [?] FnClauses
ObjectValParam	::=	(ObjectParams [?])

ObjectParams	::=	(ObjectParam ,) [*] ObjectKeyword (, ObjectKeyword) [*] ObjectParam (, ObjectParam) [*]
ObjectKeyword	::=	ObjectParam = Expr
ObjectParam	::=	FldMods [?] Param transient Param
GolnAnObject	::=	GoFrontInAnObject GoBackInAnObject [?] GoBackInAnObject
GoFrontInAnObject	::=	GoesFrontInAnObject ⁺
GoesFrontInAnObject	::=	FldDecl GetterSetterDef
GoBackInAnObject	::=	GoesBackInAnObject ⁺
GoesBackInAnObject	::=	MdDef
AbsObjectDecl	::=	ObjectHeader AbsGolnAnObject [?] end
AbsGolnAnObject	::=	AbsGoFrontInAnObject AbsGoBackInAnObject [?] AbsGoBackInAnObject
AbsGoFrontInAnObject	::=	AbsGoesFrontInAnObject ⁺
AbsGoesFrontInAnObject	::=	ApiFldDecl AbsGetterSetterDecl
AbsGoBackInAnObject	::=	AbsGoesBackInAnObject ⁺
AbsGoesBackInAnObject	::=	AbsMdDecl AbsCoercion

A.5 Variable Declaration

VarDecl	::=	VarWTypes InitVal VarWoTypes = Expr VarWoTypes : TypeRef ... InitVal VarWoTypes : SimpleTupleType InitVal
VarWTypes	::=	VarWType (VarWType (, VarWType) ⁺)
VarWType	::=	VarMods [?] Id IsType
VarWoTypes	::=	VarWoType (VarWoType (, VarWoType) ⁺)
VarWoType	::=	VarMods [?] Id
InitVal	::=	(= :=) Expr
AbsVarDecl	::=	VarWTypes VarWoTypes : TypeRef ... VarWoTypes : SimpleTupleType

A.6 Function Declaration

FnDecl	::=	FnDef AbsFnDecl
FnDef	::=	FnMods [?] FnHeaderFront FnHeaderClause = Expr
AbsFnDecl	::=	FnMods [?] FnHeaderFront FnHeaderClause Name : ArrowType

FnHeaderFront ::= Id StaticParams[?] ValParam
 | OpHeaderFront

A.7 Headers

Extends ::= extends TraitTypes
 Excludes ::= excludes TraitTypes
 Comprises ::= comprises ComprisingTypes
 TraitTypes ::= TraitType
 | { TraitTypeList }
 TraitTypeList ::= TraitType (, TraitType)^{*}
 ComprisingTypes ::= TraitType
 | { ComprisingTypeList }
 ComprisingTypeList ::= . . .
 | TraitType (, TraitType)^{*} (, . . .)[?]
 FnHeaderClause ::= IsType[?] FnClauses
 FnClauses ::= Throws[?]
 Throws ::= throws MayTraitTypes
 MayTraitTypes ::= { }
 | TraitTypes
 CoercionClauses ::= Throws[?]
 UniversalMod ::= private
 TraitMod ::= value
 | UniversalMod
 TraitMods ::= TraitMod⁺
 ObjectMods ::= TraitMods
 FnMod ::= LocalFnMod
 | UniversalMod
 FnMods ::= FnMod⁺
 VarMod ::= var
 | UniversalMod
 VarMods ::= VarMod⁺
 AbsFldMod ::= hidden | settable | UniversalMod
 AbsFldMods ::= AbsFldMod⁺
 FldMod ::= var
 | AbsFldMod
 FldMods ::= FldMod⁺
 ApiFldMod ::= hidden | settable | UniversalMod
 ApiFldMods ::= ApiFldMod⁺
 LocalFnMod ::= atomic
 LocalFnMods ::= LocalFnMod⁺
 StaticParams ::= [StaticParamList]
 StaticParamList ::= StaticParam (, StaticParam)^{*}
 StaticParam ::= Id Extends[?]
 | nat Id
 | int Id

| bool Id
| opr Op
| ident Id

A.8 Parameters

ValParam ::= BindId
| (Params[?])
Params ::= (Param ,)^{*} Keyword (, Keyword)^{*}
| (Param ,)^{*}
| Param (, Param)^{*}
Keyword ::= Param = Expr
PlainParam ::= BindId IsType[?]
| TypeRef
Param ::= PlainParam
OpHeaderFront ::= opr StaticParams[?] (LeftEncloser | Encloser) Params (RightEncloser | Encloser) (:=
| opr StaticParams[?] ValParam Op
| opr (Op | Encloser) StaticParams[?] ValParam
SubscriptAssignParam ::= Param

A.9 Method Declaration

MdDecl ::= MdDef
| AbsMdDecl
MdDef ::= FnMods[?] MdHeaderFront FnHeaderClause = Expr
| Coercion
AbsMdDecl ::= abstract ? FnMods[?] MdHeaderFront FnHeaderClause
MdHeaderFront ::= (Receiver .)[?] Id StaticParams[?] MdValParam
| OpHeaderFront
Receiver ::= Id
| self
GetterSetterDecl ::= GetterSetterDef
| AbsGetterSetterDecl
GetterSetterDef ::= FnMods[?] GetterSetterMod MdHeaderFront FnHeaderClause = Expr
GetterSetterMod ::= getter | setter
AbsGetterSetterDecl ::= abstract ? FnMods[?] GetterSetterMod MdHeaderFront FnHeaderClause
Coercion ::= *widening*[?] *coercion* StaticParams[?] (Id IsType) CoercionClauses = Expr
AbsCoercion ::= *widening*[?] *coercion* StaticParams[?] (Id IsType) CoercionClauses

A.10 Method Parameters

MdValParam ::= (MdParams[?])
MdParams ::= (MdParam ,)^{*} MdKeyword (, MdKeyword)^{*}
| (MdParam ,)^{*}
| MdParam (, MdParam)^{*}

MdKeyword	::=	MdParam = Expr
MdParam	::=	Param
		self

A.11 Field Declarations

FldDecl	::=	FldWTypes InitVal
		FldWoTypes = Expr
		FldWoTypes : TypeRef ... InitVal
		FldWoTypes : SimpleTupleType InitVal
FldWTypes	::=	FldWType
		(FldWType (, FldWType) ⁺)
FldWType	::=	FldMods [?] Id IsType
FldWoTypes	::=	FldWoType
		(FldWoType (, FldWoType) ⁺)
FldWoType	::=	FldMods [?] Id

A.12 Abstract Filed Declaration

AbsFldDecl	::=	AbsFldWTypes
		AbsFldWoTypes : TypeRef ...
		AbsFldWoTypes : SimpleTupleType
AbsFldWTypes	::=	AbsFldWType
		(AbsFldWType (, AbsFldWType) ⁺)
AbsFldWType	::=	AbsFldMods [?] Id IsType
AbsFldWoTypes	::=	AbsFldWoType
		(AbsFldWoType (, AbsFldWoType) ⁺)
AbsFldWoType	::=	AbsFldMods [?] Id
ApiFldDecl	::=	ApiFldMods [?] Id IsType

A.13 Expressions

Expr	::=	AssignLefts AssignOp Expr
		OpExpr
		DelimitedExpr
		FlowExpr
		fn ValParam IsType [?] Throws [?] ⇒ Expr
		Expr as TypeRef
		Expr asif TypeRef
AssignLefts	::=	[(AssignLeft (, AssignLeft) [*])
		AssignLeft
AssignLeft	::=	SubscriptExpr
		FieldSelection
		BindId
SubscriptExpr	::=	Primary [ExprList [?]]

FieldSelection	::=	Primary . Id
OpExpr	::=	EncloserOp OpExpr [?] EncloserOp [?] OpExpr EncloserOp OpExpr [?] Primary
EncloserOp	::=	Encloser Op
Primary	::=	Comprehension Id [StaticArgList] BaseExpr LeftEncloser ExprList [?] RightEncloser Primary [ExprList [?]] Primary . Id ([StaticArgList]) [?] TupleExpr Primary . Id ([StaticArgList]) [?] () Primary . Id Primary ^ BaseExpr Primary ExponentOp Primary TupleExpr Primary () Primary Primary TraitType . coercion ([StaticArgList]) [?] (Expr)
FlowExpr	::=	exit Id [?] (with Expr) [?] Accumulator ([GeneratorList]) [?] Expr atomic AtomicBack tryatomic AtomicBack throw Expr
AtomicBack	::=	AssignLefts AssignOp Expr OpExpr DelimitedExpr
GeneratorList	::=	Generator (, Generator) [*]
Generator	::=	GeneratorBinding Expr
GeneratorBinding	::=	Id ← Expr (Id , IdList) ← Expr

A.14 Expressions Enclosed by Keywords

DelimitedExpr	::=	TupleExpr ObjectExpr DoExpr LabelExpr WhileExpr ForExpr IfExpr CaseExpr TypecaseExpr TryExpr
TupleExpr	::=	((Expr ,) [*] (Expr ... ,) [?] Binding (, Binding) [*]) NoKeyTuple
NoKeyTuple	::=	((Expr ,) [*] Expr ...)

	((Expr ,) * Expr)
ObjectExpr	::= object Extends? GoInAnObject end
DoExpr	::= (DoFront also) * DoFront end
DoFront	::= (at Expr)? atomic? do Block?
LabelExpr	::= label Id Block end Id
WhileExpr	::= while Generator DoExpr
ForExpr	::= for GeneratorList DoFront end
IfExpr	::= if Generator then Block ElifClause* ElseClause? end [(if Generator then Block ElifClause* ElseClause end?)]
ElifClause	::= elif Expr then Block
ElseClause	::= else Block
CaseExpr	::= case Expr Op? of CaseClauses CaseElseClause? end
CaseClauses	::= CaseClause ⁺
CaseClause	::= Expr ⇒ Block
CaseElseClause	::= else ⇒ Block
TypecaseExpr	::= typecase TypecaseBindings of TypecaseClauses CaseElseClause? end
TypecaseBindings	::= (BindingList) Binding Id
BindingList	::= Binding (, Binding)*
Binding	::= BindId = Expr
TypecaseClauses	::= TypecaseClause ⁺
TypecaseClause	::= TypecaseTypeRefs ⇒ Block
TypecaseTypeRefs	::= (TypeRefList) TypeRef
TryExpr	::= try Block Catch? (finally Block)? end
Catch	::= catch Id CatchClauses
CatchClauses	::= CatchClause ⁺
CatchClause	::= TraitType ⇒ Block
Comprehension	::= { Expr GeneratorList } { Entry GeneratorList } < Expr GeneratorList > [ArrayComprehensionClause ⁺]
Entry	::= Expr ↦ Expr
ArrayComprehensionLeft	::= IdOrInt ↦ Expr (IdOrInt , IdOrIntList) ↦ Expr
ArrayComprehensionClause	::= ArrayComprehensionLeft GeneratorList
IdOrInt	::= Id IntLiteral
IdOrIntList	::= IdOrInt (, IdOrInt)*
BaseExpr	::= NoKeyTuple Literal Id self
ExprList	::= Expr (, Expr)*

A.15 Local Declarations

Block	::=	BlockElement ⁺
BlockElement	::=	LocalVarFnDecl Expr (, GeneratorList) [?]
LocalVarFnDecl	::=	LocalFnDecl ⁺ LocalVarDecl
LocalFnDecl	::=	LocalFnMods [?] FnHeaderFront FnHeaderClause = Expr
LocalVarDecl	::=	LocalVarWTypes InitVal LocalVarWTypes LocalVarWoTypes = Expr LocalVarWoTypes : TypeRef ... InitVal [?] LocalVarWoTypes : SimpleTupleType InitVal [?]
LocalVarWTypes	::=	LocalVarWType (LocalVarWType (, LocalVarWType) ⁺)
LocalVarWType	::=	var ? Id IsType
LocalVarWoTypes	::=	LocalVarWoType (LocalVarWoType (, LocalVarWoType) ⁺)
LocalVarWoType	::=	var ? Id

A.16 Literals

Literal	::=	() BooleanLiteral CharacterLiteral StringLiteral NumericLiteral
RectElements	::=	Expr MultiDimCons [*]
MultiDimCons	::=	RectSeparator Expr
RectSeparator	::=	; + Whitespace

A.17 Types

IsType	::=	: TypeRef
TypeRef	::=	TraitType ArrowType TupleType (TypeRef ?)
TraitType	::=	DottedId ([StaticArgList]) [?] { TypeRef \mapsto TypeRef } < TypeRef > TypeRef [ArraySize [?]] TypeRef ^ IntLiteral TypeRef ^ (ExtentRange (\times ExtentRange) [*])
ArrowType	::=	TypeRef \rightarrow TypeRef Throws [?]

TupleType	::=	((TypeRef ,) [*] (TypeRef ... ,) [?] KeywordType (, KeywordType) [*])
		((TypeRef ,) [*] TypeRef ...)
		SimpleTupleType
KeywordType	::=	Id = TypeRef
SimpleTupleType	::=	(TypeRef , TypeRefList)
TypeRefList	::=	TypeRef (, TypeRef) [*]
StaticArgList	::=	StaticArg (, StaticArg) [*]
StaticArg	::=	Op
		TypeRef
		(StaticArg)
ArraySize	::=	ExtentRange (, ExtentRange) [*]
ExtentRange	::=	StaticArg [?] # StaticArg [?]
		StaticArg [?] : StaticArg [?]
		StaticArg

A.18 Symbols and Operators

AssignOp	::=	:=
		Op =
Accumulator	::=	Σ Π BIG Op

A.19 Identifiers

IdList	::=	Id (, Id) [*]
Name	::=	Id
		opr Op
DottedId	::=	Id (. Id) [*]
BindId	::=	Id
		-

Bibliography

- [1] Robert D. Blumofe and Charles E. Leiserson. Scheduling multithreaded computations by work stealing. *Journal of the ACM*, 46(5):720–748, September 1999.